

Vinayak Magdum

vinumagdum114@gmail.com | +91-7058726340 | Pune,India | linkedin.com/in/vinayak-magdum62002

SUMMARY

AI/ML Engineer with an MSc in Data Science, specializing in LLM application engineering, NLP, and retrieval-augmented systems. Built and shipped a suite of AI systems at Shework covering candidate sourcing, JD generation, resume/JD parsing, candidate fit scoring, interview question generation, and answer evaluation. Experienced in building reliable LLM applications with structured outputs, schema validation, retry/backoff strategies, and graceful degradation. Strong background in document intelligence, prompt engineering, and modular AI system design.

EDUCATION

MSc Data Science and Analytics <i>Cardiff University</i>	09/2023 – 09/2024
BSc Computer Science <i>MIT-WPU</i>	07/2020 – 06/2023

PROFESSIONAL EXPERIENCE

AI/ML Engineer <i>Shework</i>	09/2025 – Present Pune,India
• Built and shipped a multi-source AI candidate sourcing engine converting JDs into ranked shortlists using LinkedIn, Naukri, recruiter-uploaded CVs, GitHub enrichment, fuzzy skill matching, deduplication, and ranking; iterated across 25 documented versions based on recruiter feedback. • Developed single and bulk LLM-based candidate fit-scoring systems with explainable recommendations, schema validation, score guardrails, concurrent processing of 3 resumes, and 8-stage performance instrumentation. • Built an interview automation pipeline using Google Gemini for JD-to-question generation and an LLM-based answer evaluation engine, with Pydantic validation, coding-question generation, retry handling, and a 26-test pytest suite using 10 recorded model-response fixtures. • Engineered a multi-format document intelligence layer supporting PDF, DOCX, TXT, and legacy DOC files using magic-byte detection, normalization, and a custom Word 97–2003 binary parser implemented in Python. • Developed a retrieval-augmented POSH compliance assistant using FAISS, MiniLM embeddings, recursive chunking, and top-k retrieval over a curated legal knowledge base. • Implemented reusable LLM reliability patterns including structured JSON outputs, schema validation, score/type normalization, retry and backoff strategies, error classification, and graceful degradation across AI workflows. • Designed modular, UI-independent processing pipelines and a stateless JSON contract for portal integration, supported by architecture documentation, versioned design plans, and formal bug-tracking records.	

SKILLS

Languages: Python, JavaScript
LLM & GenAI: Groq API, Google Gemini, Prompt Engineering, Structured Outputs, JSON Mode, LLM Output Validation, Schema Enforcement, Prompt Versioning
RAG & NLP: RAG, FAISS, sentence-transformers, MiniLM, Embeddings, Semantic Search, Text Chunking & Retrieval, NLP, Regex, RapidFuzz, Entity Resolution, Deduplication
Document AI: PDF, DOCX, DOC & TXT Processing, pdfplumber, PyPDF2, python-docx, olefile, Magic-byte File Detection, Text Normalization
Backend & Engineering: Pydantic, FastAPI Service-layer Development, REST API Integration, ThreadPoolExecutor, Concurrent Processing, Retry & Backoff, Error Classification, Graceful Degradation, pytest, structlog, Git
UI & Data: Streamlit, Web Speech API, openpyxl, Excel Reporting Automation

PROJECTS

POSH RAG <i>POSH Compliance RAG Assistant Python, FAISS, sentence-transformers, Groq</i>
Built a retrieval-augmented compliance assistant using MiniLM embeddings, recursive chunking, FAISS top-k semantic retrieval, and a curated legal knowledge base to ground generated responses in retrieved source content.